

Operating System

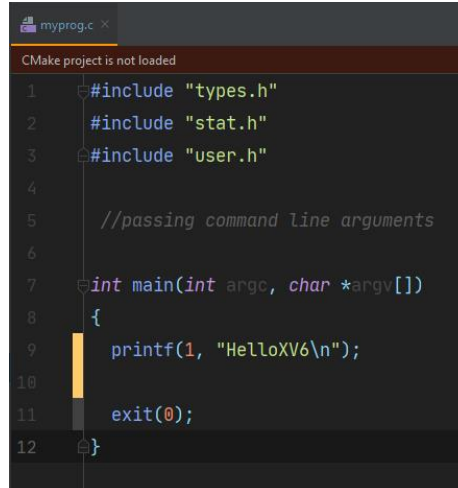
Assignment 1

Eran Aflalo 209343722

Omer Luxembourg 205500390

Task 1: Warm up ("HelloXV6"):

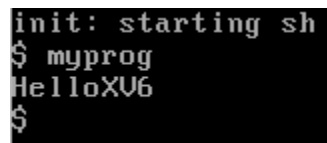
This part of the assignment is aimed at getting you started. It includes a small change in xv6 shell. Note that in terms of writing code, the current xv6 implementation is limited: it does not support system calls you may use when writing on Linux and its standard library is quite limited.



```
myprog.c x
CMake project is not loaded
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 //passing command line arguments
6
7 int main(int argc, char *argv[])
8 {
9     printf(1, "HelloXV6\n");
10
11     exit(0);
12 }
```

We created the following C program and add it to the Makefile – added myprog.c to ‘EXTRA’ section and _myprog to ‘UPROG’ section.

Results below:



```
init: starting sh
$ myprog
HelloXV6
$
```

Task 2: Support the PATH environment variable:

In this task we edit *sh.c* to maintain an array of 10 paths which are variable length up to 100 chars, that will be the environment paths as we know from Linux.

First, we define a global array for the path:

```
char env_path[10][100];
```

Next we added to the 'main' function the following code:

```
184 // Read and run input commands.
185 while (getcmd(buf, sizeof(buf)) >= 0) {
186     if (buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' ') {
187         // Chdir must be called by the parent, not the child.
188         buf[strlen(buf) - 1] = 0; // chop \n
189         if (chdir(buf + 3) < 0)
190             printf(2, "cannot cd %s\n", buf + 3);
191         continue;
192     }
193     // -- start edit 9/11 --
194     // set PATH ./bin
195     if (buf[0] == 's' && buf[1] == 'e' && buf[2] == 't' && buf[3] == ' ' &&
196         buf[4] == 'P' && buf[5] == 'A' && buf[6] == 'T' && buf[7] == 'H' && buf[8] == ' ') {
197         // Set must be called by the parent, not the child.
198         buf[strlen(buf) - 1] = 0; // chop \n
199         char * paths_line = buf + 9; // now its only the paths
200         int paths_len = strlen(paths_line);
201         for (int i = 0; i < 10; i++) {
202             if (paths_len != 0) {
203                 int s_end = 0;
204                 while ((paths_line[s_end] != ':') & (strlen(paths_line) > s_end)) {
205                     env_path[i][s_end] = paths_line[s_end];
206                     s_end++;
207                 }
208                 if (env_path[i][s_end-1] == '/') {
209                     env_path[i][s_end] = '\0';
210                 }
211                 else{ // if no '/' add it to the path
212                     env_path[i][s_end] = '/';
213                     env_path[i][s_end+1] = '\0';
214                 }
215                 paths_line = paths_line + s_end + 1; // it's a buffer - accumulate it by s_end+1
216                 paths_len = paths_len - s_end;
217
218                 // printf(1, "%d is: ", i);
219                 // printf(1, "paths- '%s' ", paths_line);
220                 // printf(1, "path is- '%s'\n", env_path[i]);
221             } else
222                 memset(env_path[i], '\0', 100);
223         }
224         continue;
225     }
226     // -- end edit 9/11 --
227     if (fork1() == 0)
228         runcmd(parsecmd(buf));
229     wait(0);
230 }
231 exit(0);
```

The code will allow us to run the command 'set PATH' and afterward the paths we want to add to the environment.

After this edition, we used the paths array to find commands to run from a current path, which are not the 'home' path ("/") or order directories, which are not the current directory – the following code is in *sh.h* in the 'runcmd' function:

```

77     case EXEC:
78         ecmd = (struct execcmd *) cmd;
79         if (ecmd -> argv[0] == 0)
80             exit(0);
81         exec(ecmd -> argv[0], ecmd -> argv);
82         // -- start edit 9/11 --
83         if(ecmd -> argv[0][0] != '/'){ // if we are NOT looking for absolute path then search path array
84             //printf(1, "We dont have absolute path... AND WE HOPE IT WILL RUN\n");
85             for(int i=0; i<10; i++){
86                 int length = strlen(env_path[i]) + strlen(ecmd -> argv[0]);
87                 char str[length + 1];
88                 // store length of env_PATH in the length variable
89                 length = 0;
90                 while (env_path[i][length] != '\0') {
91                     str[length] = env_path[i][length];
92                     length ++;
93                 }
94                 str[length] = '/';
95                 // concatenate argv[0] to str
96                 for (int j = 0; ecmd -> argv[0][j] != '\0'; j++, length++) {
97                     str[length] = ecmd -> argv[0][j];
98                 }
99                 // terminating the str string
100                str[length] = '\0';
101                exec(str, ecmd -> argv);
102            }
103        }
104        //else
105        //printf(1, "We did absolute path, BUT IT FAILED RUNNING...\n");
106        // -- end edit 9/11 --
107        printf(2, "exec %s failed\n", ecmd -> argv[0]);
108        break:

```

Results below:

```

init: starting sh
$ mkdir folder1
$ cd folder1
$ ls
exec ls failed
$ set PATH /:
$ ls
.          1 22 32
..         1 1 512
$ _

```

Before setting the PATH, we can't use the *ls* command, because it is in the *home* directory. After we added the *home* directory, which is */*, we could run the command *ls*.

* note: if we run */ls* for the directory */folder1*, the command will work, as requested.

Task 3: Extend Functionality of XV6

3.1 We added user space program called *tee* as requested.

We created a C program and added it to the Makefile. The whole code is in *tee.c*.

The first section is for a single file:

```
9 int main(int argc, char *argv[])
10 {
11     if(argc == 2){
12         char buf_read[512]; // creates a buffer
13         if (open(argv[1], 0)>=0) // a file already exists
14             unlink(argv[1]); // delete the file if exists...
15         int fd_txt1 = open(argv[1], O_WRONLY | O_CREATE); // create new fd if not exists, to write only
16         while(1){
17             int n = read(0, buf_read, sizeof(buf_read)); // return number of bytes
18             if(n==0)
19                 break; // end of file
20             else if(n<0) { // error in read
21                 printf(2, "read error\n");
22                 break;
23             }
24             else if(n>= 2 && buf_read[n-2]==3)
25                 break; // ctrl + c
26             else { // write to file and console
27                 if (write(fd_txt1, buf_read, n) != n) { // written bytes not equal to n = error!
28                     printf(2, "write error\n");
29                     break;
30                 }
31                 if (write(1, buf_read, n) != n) { // written bytes not equal to n = error!
32                     printf(2, "write to console error\n");
33                     break;
34                 }
35             }
36         }
37         close(fd_txt1); // closing file descriptor
```

The results of a single file:

Echo after each Enter

This is CTRL+C and then Enter

```
init: starting sh
$ tee tmp.txt
a
a
b
b
c
c
c
^C
$ cat tmp.txt
a
b
c
c
$ _
```

For two files we created the following code:

```
38 }else if(argc == 3){
39     char buf_read[512]; // creates a buffer
40     int fd_txt1 = open(argv[1], O_RDONLY);
41     if (open(argv[2], 0)>=0) // a file already exists
42         unlink(argv[2]); // delete the file if exists...
43     int fd_txt2 = open(argv[2], O_WRONLY | O_CREATE); // create new fd if not exists, to write only
44     while(1){
45         int n = read(fd_txt1, buf_read, sizeof(buf_read)); // return number of bytes
46         if(n==0)
47             break; // end of file
48         else if(n<0) { // error in read
49             printf(2, "read error\n");
50             break;
51         }
52         else { // write to file and console
53             if (write(fd_txt2, buf_read, n) != n) { // written bytes not equal to n = error!
54                 printf(2, "write error\n");
55                 break;
56             }
57         }
58     }
59     close(fd_txt1); // closing file descriptor
60     close(fd_txt2); // closing file descriptor
61 }else
62     printf(2, "cannot tee, given %d arguments\n", argc);
63 exit(0);
64 }
```

The results are as follows:

```
console      3 21 0
folder1      1 22 32
tmp.txt      2 23 6
$ cat tmp.txt
a
b
c
$ tee tmp.txt out.txt
$ cat out.txt
a
b
c
$ _
```

As we can see the content of *tmp.txt* is in *out.txt* now.

3.2 *getpinfo* is a system call that print a list of currently running processes.

To implement it, we created a user space program and a kernel system call as follows:

User space program only calls the system call –

```
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  #include "fcntl.h"
6
7  //passing command line arguments
8
9  int main(int argc, char *argv[])
10 {
11     getpinfo();
12     exit(0);
13 }
```

To create the kernel system call we added the following files with appropriate function calls and macros:

user.h *ysys.S* *syscall.c* *syscall.h*

In addition, we created the following function to *proc.c* file:

```
545 void
546 getpinfo(void)
547 {
548     struct proc *p;
549     int i = 0;
550
551     acquire(&ptable.lock);
552     cprintf("<Number of Row>\t\t<Process ID>\n");
553     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
554         if (p->state != UNUSED) {
555             cprintf("\t%d\t\t\t%d\n", i, p->pid);
556             i++;
557         }
558     }
559     release(&ptable.lock);
560 }
```

Results:

```
init: starting sh
$ getpinfo
<Number of Row>          <Process ID>
                        0                1
                        1                2
                        2                3
$
```

* Note that the QEMU terminal does not identify "\t" as a regular tab.

Task 4: Wait and exit system calls

4.1 We updated the process structure with a status field in *proc.h* (and updated *user.h*, *defs.h*, *sysrpoc.c*, *proc.c*). In addition, we edited *exit* function (in *proc.c*) signature to receive an *int* status.

proch.h – added a field

```
37 // Per-process state
38 struct proc {
39     uint sz;           // Size of process memory (bytes)
40     pde_t* pgdir;      // Page table
41     char *kstack;      // Bottom of kernel stack for this process
42     enum procstate state; // Process state
43     int pid;           // Process ID
44     struct proc *parent; // Parent process
45     struct trapframe *tf; // Trap frame for current syscall
46     struct context *context; // Switch() here to run process
47     void *chan;        // If non-zero, sleeping on chan
48     int killed;        // If non-zero, have been killed
49     struct file *ofile[NOFILE]; // Open files
50     struct inode *cwd; // Current directory
51     char name[16];     // Process name (debugging)
52     // -- start edit 13/11 --
53     int status;        // added status as int
54     // -- end edit 13/11 --
55 };
```

proc.c – exit function

```
227 void
228 exit(int status) // 13/11 changed from void to int
229 {
230     struct proc *curproc = myproc();
231     struct proc *p;
232     int fd;
```

...

```
264 // -- start edit 13/11 --
265 curproc->status = status; // assigning the status to the current process
266 // -- end edit 13/11 --
267
268 // Jump into the scheduler, never to return.
269 curproc->state = ZOMBIE;
270 sched();
271 panic("zombie exit");
272 }
```

sysproc.c –

```
17 sys_exit(void)
18 {
19     // -- start edit 13/11 --
20     int status;
21     if(argint(0, &status) < 0)
22         return -1;
23     exit(status);
24     // -- end edit 13/11 --
25     //exit();
26     return 0; // not reached
27 }
```

After changing *exit* function, we changed “*exit()*” to “*exit(0)*” where needed.

4.2 Updating the wait system call. We changed *wait* system call to resolve with the status given in *exit*.

Same as *exit* we updated the files:

proc.c -

```
275 // Return -1 if this process has no children.
276 int
277 wait(int *status) // 13/11 changed from void
278 {
279     struct proc *p;
280     int havekids, pid;
281     struct proc *curproc = myproc();
282
283     acquire(&table.lock);
284     for(;;){
285         // Scan through table looking for exited children.
286         havekids = 0;
287         for(p = table.proc; p < &table.proc[NPROC]; p++){
288             if(p->parent != curproc)
289                 continue;
290             havekids = 1;
291             if(p->state == ZOMBIE){
292                 // Found one.
293                 pid = p->pid;
294                 kfree(p->kstack);
295                 p->kstack = 0;
296                 freevm(p->pgdir);
297                 p->pid = 0;
298                 p->parent = 0;
299                 p->name[0] = 0;
300                 p->killed = 0;
301                 p->state = UNUSED;
302                 // -- start edit 13/11 --
303                 if (status != 0) // if status is NULL we cannot assign value to it
304                     *status = p->status;
305                 // -- end edit 13/11 --
306                 release(&table.lock);
307                 return pid;
308             }
309         }
310     }
```

The status in this function is a pointer to status field in 'p' process.

sysproc.c -

```
28
29 int
30 sys_wait(void)
31 {
32     // -- start edit 13/11 --
33     int *status;
34     if(argptr(0, (void*)&status, sizeof(int*)) < 0)
35         return -1;
36     return wait(status);
37     // -- end edit 13/11 --
38     //return wait();
39 }
```

We created a user space program to test the functionality of the status with wait and exit. The following results show the status value that got from the function by calling "*exit(57)*".

This is the user space program “wait_test”:

```
9  int main(int argc, char *argv[])
10  {
11      int pid = fork();
12      if(pid == 0){
13          printf(1, "I'm the child!\n");
14          exit(57);
15      } else if(pid < 0){
16          printf(1, "This is fork failed\n");
17          exit(0);
18      }
19      int status;
20      wait(&status);
21      printf(1, "I'm the parent, and got child status: %d\n", status);
22      exit(0);
23  }
```

The terminal results:

```
init: starting sh
$ wait_test
I'm the child!
I'm the parent, and got child status: 57
$
_
```

4.3 Updating the wait system call – we adjust the “exec.c” file as requested.

exitf() – will create the function for the stack.

```
14  void
15  exitf()
16  {
17      asm volatile ("push %eax;");
18      asm volatile ("pop 4(%esp);");
19      asm volatile ("mov %0 , %%eax" : : "i" (SYS_exit));
20      asm volatile ("int %0" : : "i" (T_SYSCALL));
21  }
```

In the exec function we updated the following:

```
23  int
24  exec(char *path, char **argv)
25  {
26      char *s, *last;
27      int i, off;
28      uint argc, sz, sp, ustack[3+MAXARG+1];
29      struct elfhdr elf;
30      struct inode *ip;
31      struct proghdr ph;
32      pde_t *pgdir, *oldpgdir;
33      struct proc *curproc = myproc();
34
35      //STEP 1 exit syscall size(byte)
36      int exitsz = (int)exec - (int)exitf;
37  }
```

```

88 //Step 2 - stack pointer update and copy page from sp to exitf
89 sp = sz - exitsz;
90 copyout(pgdir, sp, exitf, exitsz);
91
92 // Push argument strings, prepare rest of stack in ustack.
93 for(argc = 0; argv[argc]; argc++) {
94     if(argc >= MAXARG)
95         goto bad;
96     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
97     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
98         goto bad;
99     ustack[3+argc] = sp;
100 }
101 ustack[3+argc] = 0;
102
103 ustack[0] = sz - exitsz; // NO fake return PC - return address of user space program
104 ustack[1] = argc;
105 ustack[2] = sp - (argc+1)*4; // argv pointer

```

To test it we emitted the “exit(0)” in the user space program “wait_test.c” that we created before:

The screenshot shows a QEMU virtual machine window. On the left, a code editor displays the source code of 'wait_test.c'. The code includes headers for 'types.h', 'stat.h', 'user.h', and 'fcntl.h'. It defines a 'main' function that forks a child process. The child process prints 'I'm the child!', calls 'exit(57)', and then the parent process prints 'I'm the parent, and got child status: 57' before calling 'exit(0)'. On the right, the QEMU console shows the boot process: SeaBIOS version 1.10.2-1ubuntu1, iPXE boot, and QEMU booting from a hard disk. The console output shows the user-space program 'wait_test' being executed, with the parent process printing 'I'm the parent, and got child status: 57' and the child process printing 'I'm the child!' before exiting with status 57.

```

1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 #include "fcntl.h"
6
7 //passing command line arguments
8
9 int main(int argc, char *argv[])
10 {
11     int pid = fork();
12     if(pid == 0){
13         printf(1, "I'm the child!\n");
14         exit(57);
15     } else if(pid < 0){
16         printf(1, "This is fork failed!\n");
17         exit(0);
18     }
19     int status;
20     wait(&status);
21     printf(1, "I'm the parent, and got child status: %d\n", status);
22     //exit(0);
23 }

```

```

QEMU - Press Ctrl-Alt to exit mouse grab
SeaBIOS (version 1.10.2-1ubuntu1)

iPXE (http://ipxe.org) 00:03.0 C980 PCI2.10 PnP PMM+1FF8DDDD+1FECDDDD C980

Booting from Hard Disk...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap star
t 58
init: starting sh
$ wait_test
I'm the child!
I'm the parent, and got child status: 57
$ ?_

```

As expected, the adjustment worked!